

Introduction to Networking and Systems Measurements

Measurement Pitfalls



Dr Andrew W. Moore

andrew.moore@cl.cam.ac.uk

Dr Noa Zilberman

noa.zilberman@cl.cam.ac.uk

Common Measurement Pitfalls

- What are the hidden assumptions?
- What did you not notice (in the system, setup,)?
- What can your tool do?
- Vantage points
- Latency pitfalls
- BW pitfalls
- Reading the results

Hidden Assumptions - Examples

- The path from A to B is the same (reverse) as the path from B to A
- There is no packet reordering
- Device throughput is the same for all packet sizes
- Test packets will experience the same effects as application's traffic
- The effect of DNS lookup is negligible
- The measurement tool has negligible overhead
- Previous work was correct

System and Setup

Did you notices that....

- There are other jobs running on the same core
- ICMP traffic is throttled by the OS
- CPU frequency scaling is enabled
- The CPU that you are using is not connected directly to the NIC
- Kernel version has been updated overnight
- The 2x40G NIC uses PCIe Gen 3 x8 (~60Gbps)
- There is a new Errata...

What can your tool do? - Examples

- SSD can write at 450MB/s
 - Don't try to write data captured at 10Gbps
- The latency for reading CPU timestamp is ~tens of cycles
 - Don't try to use it to measure cache access time
- DAG resolution is 4ns
 - Don't try to measure the propagation delay through 1m fibre
- OSNT can only capture at low rate
 - Don't try to measure latency of 10Gbps flows

Latency Pitfalls - Examples


- What is the definition of “latency”?
 - Propagation delay? Inter packet gap? Round trip time? Flow completion time?
- How was the latency measured?
 - Start of packet to start of packet? Start of packet to end of packet?
 - Single packet? Packet-pair? Packet-train?
- Where was the timestamp taken?
 - ...and how did it affect the measurement?
- Resolution, precision and accuracy...

Bandwidth Pitfalls - Examples

- What is the definition of “bandwidth”?
 - Link capacity? Average throughput? Peak throughput?
- Controllability
 - Packet size? Protocol? QoS?
- What was the status of the network?
- Net neutrality?
- Did you pass through the bottlenecks?
- Resolution, precision and accuracy...

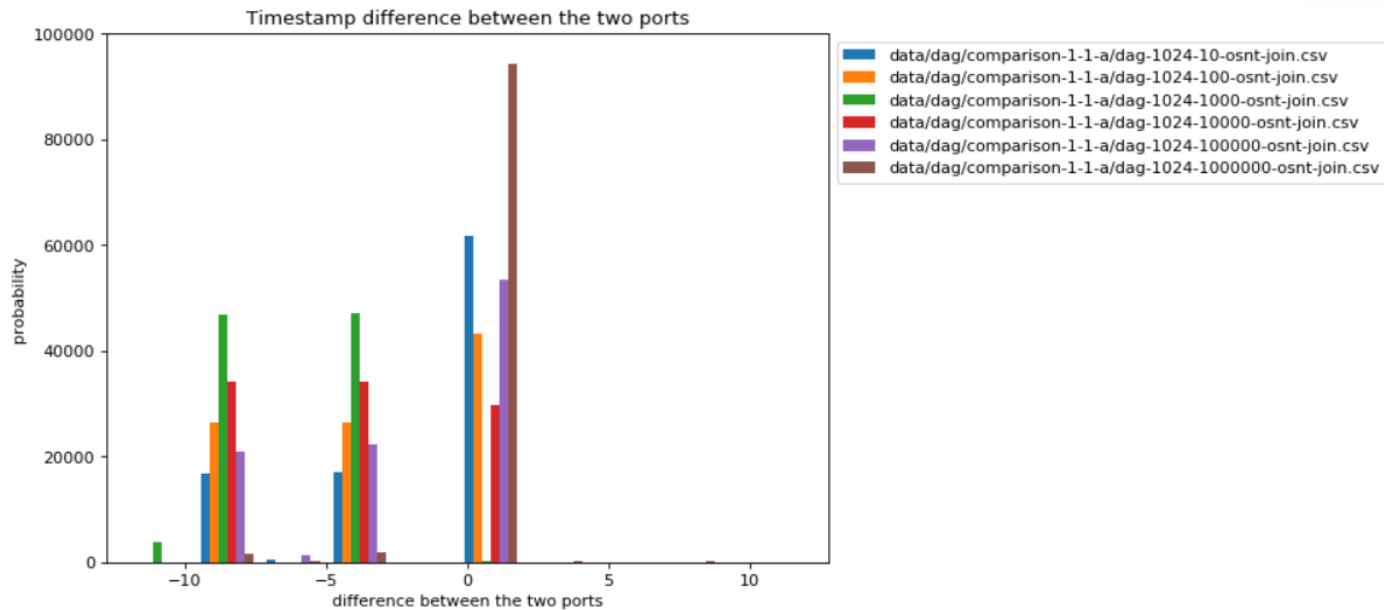
Example: Timestamp difference between ports

- Recall Lab 2, experiment 2.1 b
- Measuring the timestamp difference between 2 ports:




Example: Timestamp difference between ports

- 100,000 packets, **1024B**
- Different Inter Packet Gaps (IPG)




Example: Timestamp difference between ports

- 100,000 packets, **64B**
- Different Inter Packet Gaps (IPG)



Example: Timestamp difference between ports

- 100,000 packets, **64B**, running 10 times
- Same Inter Packet Gap (IPG)



Example: Switch Throughput

- The reported iperf result for a NetFPGA reference switch is 9.4Gbps
- User complaint: I see only 8.9Gbps and packet drop in the switch

Connecting to host 10.0.0.13, port 5201

[4] local 10.0.0.12 port 54764 connected to 10.0.0.13 port 5201

[ID]	Interval		Transfer	Bandwidth	Retr	Cwnd
[4]	0.00-1.00	sec	1.02 GBytes	8.76 Gbits/sec	74	313 KBytes
[4]	1.00-2.00	sec	1.03 GBytes	8.86 Gbits/sec	34	198 KBytes
[4]	2.00-3.00	sec	1.03 GBytes	8.87 Gbits/sec	34	281 KBytes
[4]	3.00-4.00	sec	1.04 GBytes	8.92 Gbits/sec	34	238 KBytes
[4]	4.00-5.00	sec	1.04 GBytes	8.93 Gbits/sec	32	208 KBytes
[4]	5.00-6.00	sec	1.04 GBytes	8.92 Gbits/sec	29	187 KBytes
[4]	6.00-7.00	sec	1.04 GBytes	8.95 Gbits/sec	27	365 KBytes
[4]	7.00-8.00	sec	1.04 GBytes	8.94 Gbits/sec	28	233 KBytes
[4]	8.00-9.00	sec	1.03 GBytes	8.88 Gbits/sec	30	420 KBytes
[4]	9.00-10.00	sec	1.04 GBytes	8.96 Gbits/sec	33	423 KBytes


[ID]	Interval		Transfer	Bandwidth	Retr	
[4]	0.00-10.00	sec	10.4 GBytes	8.90 Gbits/sec	355	sender
[4]	0.00-10.00	sec	10.4 GBytes	8.90 Gbits/sec		receiver

Example: Switch Throughput

- Debug: Have you tried changing rx-usec?
 - User: no more packet drop in the switch!
 - ...but bandwidth is down to 7.5Gbps...
-
- New insight: NIC used on reference setup (Solarflare) is different than the NIC used by user (Intel)
 - (skipping a few steps forward)


Example: Switch Throughput

- Switch throughput over time (10ms sampling resolution)




Example: Switch Throughput

- Switch throughput over time (100 μ s sampling resolution)



Example: Switch Throughput

- What else is different?



Example: TSC Access


- Goals:
 - Evaluate the accuracy & precision of time-taking using CPU time stamp counter (TSC)
- Methodology:
 - Read TSC twice
 - Measure the time-gap between the two consecutive reads
- Results:
 - Min/Median/99.9%: 9ns/10ns/11ns

Example: TSC Access

```
1  while (!done)
2  {
3      //Read TSC twice, one immediately after the other
4      do_rdtscp(tsc, cpu);
5      do_rdtscp(tsc2, cpu2);
6      //If the gap between the two reads is above a
          certain threshold, save it
7      if ((tsc2 - tsc > threshold) && (cpu == cpu2))
8          buffer[samples++] = tsc2 - tsc;
9  }
```

Example: TSC Access

What happens over time?






Example: TSC Access

- Source data:

X ≤	User space Events
10	91428291492
11	404700
12	268521
22	268291
120	267465
1097	10768
10869	1

X ≤	Kernel Events
9	11117819727
10	3973891503
49	287
53	201
98	90
1155	86
1184	85
1241	77
1982	1

Example: TSC Access



Example: Topology Measurements

- Goal:

- Build a map of network connectivity that assigns IP addresses to locations


- Method:

- Simple option: name resolution
 - 4.69.166.1 \Rightarrow ae-119-3505.edge4.London1.Level3.net
 - But many times information is missing, not indicative or is inaccurate
- Better option: use geolocation services
 - Most services claim to be over 99% accurate

Example: Topology Measurements

Building a map of the network:

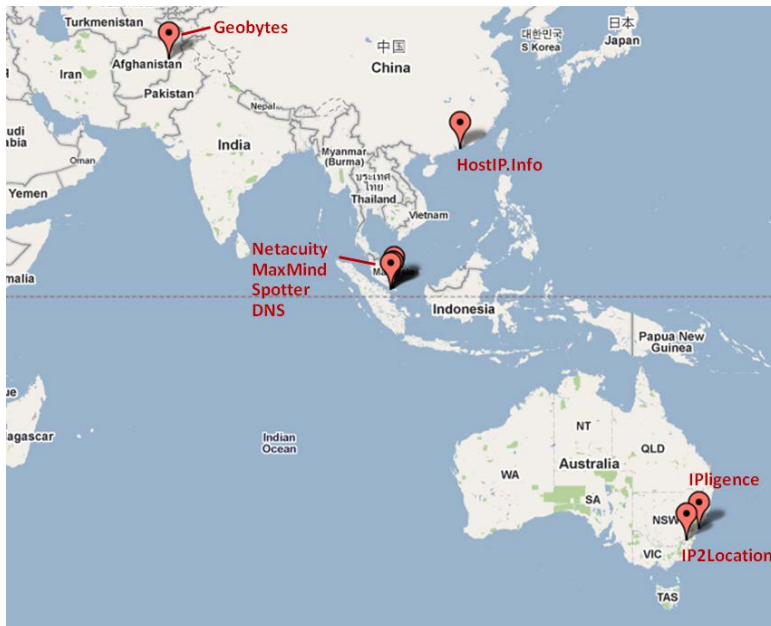
- Measurements for connectivity
- Geolocation databases for location




Example: Topology Measurements

What is your ground truth?

- Gelolocation databases are over 99% accurate!



Verizon/MCI/UUNET (ASN 703)
10-nodes PoP



Heatmap – Median distance between databases
(2011)

Validation

- Measurements need to be validated
- Don't make assertions!
- Use ground truth (where available)
- Compare different tools and methodologies
- Do the results make sense?
 - RTT can't be faster than traveling at the speed of light...
- Have I mentioned validation?



Final Report - Recommendations

- Include all figures within the report
 - Use proper scale, adapt the template if need be
- Make sure that your environment does not affect the results
- Do not make assertions
 - Support your claims through experimentations
- Discuss your results in depth:
 - Compare and contrast results gained through different vantage points, using different tools, on different platforms etc
 - Provide side-by-side comparisons
 - Use the questions in the handouts as guiding examples
- Use the right terminology (accuracy, precision, resolution)
- Correct typos and grammar mistakes
- Make sure not to run out of budget
- Follow the instructions in the handout

Course Summary

- This course covered measurements tools and measurement techniques
- But also “why out most basic assumptions are wrong”, “graphs lie”, “what you don’t know about your system”, ...
- Remember:
 - Constant vigilance
 - Look at the data, best-practice, think.
- Applies to all types of measurements

